

A Software Cryptoprocessor for Commodity Servers

Authors omitted for blind review

ABSTRACT

Defending computers from unauthorized physical access and malicious hardware devices has proved extremely challenging. The problem is exacerbated in cloud-computing environments, where users lack physical control over the server hardware that executes their workloads. Moreover, cloud service providers can be compelled by government agencies to provide physical access to servers.

We introduce *vCage*, a software-based cryptoprocessor system that employs cryptographic techniques to provide confidentiality for unmodified workloads. Code and data are stored as cleartext only within the processor cache, but remain encrypted in main memory. We implemented vCage in a commercial product for commodity server platforms based on Intel x86 processors and industry-standard TPMs.

We present the design and implementation of vCage in Linux, and show how it virtualizes physical security by protecting data-in-use transparently for KVM virtual machines. Quantitative experiments demonstrate that vCage effectively safeguards privacy, while achieving acceptable performance for a wide range of workloads.

Categories and Subject Descriptors

D.2.6 [Operating Systems]: Security and Protection

General Terms

Security

Keywords

Memory Protection, Processor Cache, Operating Systems, Hypervisors, Virtual Machines

1. INTRODUCTION

Introduction goes here...

2. PHYSICAL ATTACKS

It is generally understood that physical access to an x86 platform can completely compromise software security. Historically, physical security controls, such as cages, cameras, and locks, have been employed to prevent or detect physical access. Yet with adoption of outsourced infrastructure and cloud computing, x86 platforms are increasingly run outside the physical control of the software owner.

This section briefly summarizes several well-known physical attack vectors against x86 platforms, including DMA, physical memory extraction, and platform malware.

2.1 Direct Memory Access

By design, x86 architectures provide direct memory access (DMA) from hardware subsystems to access main memory independently of the CPU. DMA access is generally used for performance, for example allowing disk, network, and graphics devices to read and write data directly to memory without incurring CPU cycles.

Yet, without proper controls, devices with DMA access can read arbitrary regions of memory and compromise system security by exposing secrets or allowing an attacker to modify running software in place. Secrets from captured memory can be extracted easily with forensics tools like Volatility [21].

Tribble [4, 7] is an early example of a device designed for exfiltrating data via DMA access, based on an off-the-shelf Intel development kit. Copilot [18] also used DMA with a PCI device for the purpose of monitoring kernel integrity. The Moux attack [19] exploited remote vulnerabilities in a standard network interface device and accessed memory via DMA. Off-the-shelf intelligent network adapters, such as those made by Cavium, are able to exfiltrate DMA-accessed memory over a network connection [9].

The IEEE 1394 Firewire interface also provides DMA by design. This led to several demonstrations of memory extraction to steal data or for forensics [5, 2, 22, 6]. The Thunderbolt interface also offers DMA access and can be exploited in a similar manner [13].

2.2 Physical Memory Extraction

While DMA may be mitigated by software-based countermeasures, leveraging a hardware IOMMU such as Intel VT-d [11], other attacks involve the physical extraction or modification of system memory. For example, memory bus analyzer devices are available and can interdict memory traffic. However, bus analyzers must be installed ahead of time, tend to be relatively expensive, and physically large.

A “cold boot” is a low-cost memory extraction attack that

can be conducted for little cost on a running system [8]. Cold booting involves literally freezing system memory modules with an aerosol freeze spray. The memory contents are preserved long enough to boot to a “scraper” image such as *bios-memimage* or *msrاندump* which can preserve the memory contents to persistent storage.

Cold booting disrupts a running system and data must be recovered immediately, before the memory module thaws. Furthermore, it does not reliably capture all memory contents, as there is some degradation over time. Conducting the attack may be further complicated by error-correcting memory which is cleared on a reset or by data scrambling for power supply noise suppression [15].

Non-volatile memory (NV-RAM), designed to persist data after a power loss, is now available in DDR3/DDR4 form-factors used by standard x86 servers. An attacker can install NV-RAM modules in a server and remove them at any moment, recovering the data at a later time. If a memory mirroring mode is configured, an attacker could remove a NV-RAM module from a running system and replace it without disrupting service.

2.3 Bootkits and Platform Malware

- APT

2.4 Practical Attack Vectors

SW: This might be unnecessary. Do we really need to explain how people get physical access?

CW: Suggest writing a few sentences, and moving to intro for this section. The NSA attacks seem particularly interesting to at least mention.

- Supply chain
- Datacenter insiders
- Previous bare-metal host tenants: Can we pwn our SoftLayer BIOS?
- Stoned bootkit in the wild
- NSA ANT

3. CACHE MANAGEMENT

Compared to a conventional system, vCage treats the processor cache like main memory, and main memory like an encrypted backing store. All cleartext data must remain resident in the processor cache for correctness; evictions could leak sensitive data to untrusted main memory. Since cache space is smaller than traditional main memory sizes by orders of magnitude, it must be managed efficiently.

Fortunately, cache sizes are growing rapidly along with advances in process technology. Modern multi-core processors integrate a large, physically-indexed last-level cache (LLC) shared among multiple cores. Currently, the Intel Ivy Bridge EX x86 processor is available with a 37.5MB on-die LLC. Unfortunately, x86 processors lack hardware support for explicit software control over cache allocation and eviction policies. Moreover, recent x86 hardware introduced complex cache indexing, which makes software cache management even more difficult.

After reviewing processor caching behavior, including memory types and cache indexing, we explain how vCage manages cache contents despite numerous challenges.

3.1 Memory Types

Processors commonly provide some method for specifying how regions of physical memory should be cached. The x86 architecture [10] supports a small set of memory type range registers (MTRRs), each of which can be programmed to control the caching method used for a contiguous region of memory, subject to various size and alignment constraints. A similar capability is provided for specifying caching attributes of individual pages via page table entries and the x86 page attribute table (PAT).

vCage uses MTRRs to prevent nearly all of untrusted main memory from being cached, setting the memory type to *write-combining (WC)*, instead of *uncacheable (UC)*. Neither memory type is cached, but WC offers significantly higher performance than UC, due to load and store buffering. However, write combining uses a weak memory ordering model which is not appropriate for general-purpose data; in most systems, it is used only for video frame buffers. vCage accesses WC memory in a controlled manner for encrypted paging. Data is copied between the cache and memory using non-temporal move instructions that utilize write-combining store buffers and streaming load buffers efficiently [10].

3.2 Traditional Cache Indexing

Memory is cached in units of *lines*; on modern x86 hardware, the cache line granularity is 64 bytes, with 64-byte alignment. The processor maps each line of memory to a single *n*-way associative *cache set* based on its physical address. For example, on the Intel Sandy Bridge x86 processor, a single last-level cache set consists of 20 lines; *i.e.*, the cache is 20-way set associative.

Main memory is normally several orders of magnitude larger than the cache, so many more than *n* lines of physical memory map to the same cache set. When a new line of memory needs to be cached in a set that is already full, these *conflicting* lines contend for the same scarce cache space. The processor implements a cache replacement policy that chooses some existing line to replace, typically evicting one of the least recently used (LRU) lines from the set to optimize for temporal locality.

Traditionally, processors have employed a simple mapping of physical addresses to cache sets, based on lower-order bits of the physical address. For example, consider an Intel Nehalem x86 processor with an 8MB 16-way set-associative cache. Each cache set contains $16 \text{ lines} \times 64\text{B} = 1\text{KB}$ of cached data; the entire cache contains 8K sets, requiring a 13-bit index. A physical address is mapped to a cache set using bits 6..18 as the cache set index, with bits 0..5 specifying the intra-line byte offset.

3.3 Page Coloring

When paging is used for address translation, the low-order bits of the physical page number are the high-order bits of the cache set index. This straightforward hardware mapping of physical addresses to cache sets has been leveraged for many years by operating systems and hypervisors, using a technique known as *page coloring*. By controlling the virtual-to-physical page mapping, system software can partition pages into disjoint sets, or *colors*, such that pages with different colors do not conflict in the cache.

In the Nehalem example above, there are 128 page colors when using conventional 4KB x86 “small” pages. Bits 0..6 of the physical page number, corresponding to physical address

bits 12..18, specify the page color – a contiguous range of sets used to cache the page contents.

Page coloring has been used in many systems to improve performance by reducing cache conflict misses [3], and to control the isolation or sharing of cache memory between software contexts [23]. Commercial hypervisors have preserved the relative page coloring used by the guest OS within a virtual machine [20]. We had originally intended to leverage page coloring to enable vCage to manage the processor cache. However, as detailed below, recent changes in processor cache indexing have complicated this approach.

3.4 Complex Cache Indexing

Recent Intel x86 processors, starting with the Sandy Bridge micro-architecture, use a ring-based interconnect between cores and a multi-banked last level cache (LLC), which is organized into multiple cache *slices*. Physical addresses are mapped to LLC cache sets within slices using “complex cache indexing” [?, 10]. The hardware that realizes this mapping implements an undocumented, proprietary hash function, and may potentially use any physical address bits as inputs. The hash function may also vary across different processor implementations and configurations.

As a result, traditional page-coloring techniques may no longer work. A small contiguous memory region can be scattered across many discontinuous sets throughout the cache. Since vCage requires strict control over the cache, we developed a reliable approach for generating fine-grained address-to-cache-set mappings on modern Intel processors.

3.5 Discovering Set Mappings

We present a programmatic method for computing complete information about processor-specific mappings from physical addresses to cache sets. Although we use Intel x86 processors, the same general technique should work on any architecture with deterministic cache indexing and accurate cache performance counters.

As input, the method takes a collection of physical addresses at cache-line granularity, typically a single contiguous address range. As output, these addresses are partitioned into cache *conflict sets*, such that addresses within a partition conflict in the cache, and addresses in different partitions do not conflict. The basic idea is to exceed the limited associativity of a single n -way associative cache set to force observable cache evictions and discover conflicts.

An empty conflict set S is allocated, and used to maintain an array of addresses observed to conflict in the LLC. The entire cache is flushed using the the x86 `WBINVD` instruction to start from a known empty state.

Loads are issued to each input memory address, one-by-one, until a hardware performance counter, programmed to monitor LLC evictions, and checked after each load, detects that a first eviction has occurred. On the Intel Sandy Bridge and Ivy Bridge processors, the uncore caching-agent (CBo) performance counters were programmed to monitor the `LLC_VICTIMS` event, filtered by the MES cache states [?]. Since lines in different slices cannot conflict, and the hardware provides per-slice counters, this procedure is run for each slice independently. Conveniently, this also eliminates any potential interference from concurrent cache activity in other slices.

The address which caused the eviction is added to S . The cache is again flushed, and loads are performed to all ad-

resses in S , to ensure that they are resident in the cache. As a result, a different conflicting input address associated with the same set will cause the next eviction – essentially rotating through a ring of conflicting lines that exceed the hardware cache associativity. Loads are issued again to the remaining set of potentially-conflicting input addresses, one-by-one, in the same order, and the address causing the next eviction is added to S .

This process is repeated until $|S| > n$, or several repetitions yield addresses already in S . Each generated partition will contain $n + 1$ physical lines for an n -way set-associative cache. To support larger input regions, once a conflict set S contains $n + 1$ addresses, one is selected, moved to an “overflow” list associated with S , and removed as an active input address. This process can be repeated until all input memory addresses have been partitioned into conflict sets.

After each conflict set is identified, each of its constituent addresses is marked as already “used” by some set. When input addresses are read one-by-one in the method outlined above, any which have already been marked as used are instead flushed from the cache using the x86 `CLFLUSH` instruction. This prevents noise due to prefetching or eviction events from other sets.

Additional steps can further improve the robustness and accuracy of this approach, including mapping the primary data structures as uncached, disabling interrupts during the computation of each conflict set, disabling prefetching, and booting in uniprocessor mode to prevent memory accesses from other cores sharing the LLC. More generally, multiple runs can be performed to resolve any discrepancies and ensure consistent results.

3.6 Experimental Results

We implemented our set-mapping method as a loadable kernel module for Linux. There were few differences in the partitions computed by separate runs; we observed a single-line difference in less than 0.2% of the partitions computed for a cache-sized region, even without using all of the noise-reduction techniques described above. The results revealed many interesting facts about the caching behavior of the Intel Sandy Bridge and Ivy Bridge processors that can be leveraged to perform page-level partitioning.

A Sandy Bridge processor configured with a 20MB LLC contains 320K 64-byte cache lines, grouped into 16K 20-way-associative cache sets. Cache sets are partitioned across 8 cache slices, each containing 2K sets; an intra-slice cache set index can be represented in 11 bits. The conflict-set data revealed that the 17 low-order physical address bits are identical within a single set. These bits encode an 11-bit intra-slice index, plus a 6-bit intra-line byte offset. At 4KB page granularity, there are 32 cache partitions, each with size 640KB, based on address bits 12..16, *i.e.*, the lower-order page-number bits. This information yields a simple mechanism for performing coarse cache partitioning that is effectively the same as traditional cache coloring.

While the intra-slice cache index can be extracted directly from a physical address, each slice has a cache set at the same index. Two pages won’t conflict in the cache if their constituent lines are mapped to different slices. The conflict-set data also revealed patterns in the hardware address-to-slice mapping. The data indicated that the 64 consecutive lines within a page are striped across different slices in one of eight regular patterns. This information yields a second

partitioning method that is distinct from traditional page coloring, based on classifying each page into one of these eight patterns. Pages with different slice patterns do not conflict in the cache, but the slice number isn't coded as a simple bit range in the physical address.

Using both methods of cache partitioning derived from the conflict-set data, it is possible to construct "2D" nested partitions – one using address ranges, where address bits 12..16 encode the "page color", and the other using the page's slice-pattern classification. The slice classification can be stored compactly in 3 bits per 4K page using a simple lookup table. With additional analysis, it should be possible to determine the exact address-to-slice hash function, either manually, or by using machine-learning techniques. Experimentally, we determined that the two partitioning methods are orthogonal, so each 640KB page-color partition can be sub-partitioned into 8 slice patterns, yielding a smaller, more flexible 80KB partitioning granularity. This is as fine-grained as possible given the hardware – 4KB pages times 20-way set associativity.

We also found that some cache sets in slices 0 and 1 appeared to have an effective associativity of only 19 ways, instead of the expected 20. We believe this is due to hardware way-partitioning [?], allowing the integrated graphics controller to claim a dedicated portion of the LLC.

3.7 Other Approaches

The x86 architecture provides additional cache operating modes that initially seemed promising for vCage cache management. A well-documented *no-fill cache mode* can be specified by setting `CR0.CD` [10]. In this mode, read hits access the cache, but read misses are prevented from updating the cache. Write hits update the cache, but only write misses and writes to shared lines update system memory. Invalidations are also allowed. Unfortunately, this mode applies not only to the LLC, but to *all* levels of the cache hierarchy, including L1 and L2. As a result, overall performance is unacceptable, as the contents of the fastest caches remain frozen, and cannot be updated to reflect changes in the hot working set.

A separate *no-evict cache mode*, typically available only to BIOS writers under NDA for system initialization, prevents cache evictions. Unfortunately, this mode has so many onerous restrictions (such as lacking support for paging), that it is unsuitable for general-purpose execution.

- CARMA
- Non-Intel architectures: ARM and AMD?

3.8 Detecting Cache Leaks

Careful cache management, leveraging both the memory types and partitioning information discussed above, should be sufficient to ensure that all cleartext data remains resident in the cache. Nevertheless, vCage programs the uncore caching-agent (CBo) hardware performance counters [?] to monitor cache evictions and lookup misses continuously, while the system is executing.

During vCage development, we detected and fixed several potential cache leaks, mostly related to I/O and execution of system-management-mode (SMM) code.

- SMM, SMM Transfer monitor
- Intel DDIO

- Video card, using external card
- Remediation for detected leaks?

4. SYSTEM ARCHITECTURE

4.1 Cache-Resident OS

CW: Promoted Cache Management to its own section. May want to do the same for other subsections, and have the system architecture be a shorter overview, also serving as a roadmap for the rest of the paper.

Entire OS stack remains resident in the processor cache.

- Necessary kernel changes / gotchas
 - Minimizing kernel size
 - * Removing unneeded drivers
 - * Memory allocator tweaks
 - * Example: `struct page`
 - Boot order tweaks
 - * Two stage initrd
 - Never failing allocation
 - Other shrinking memory footprint

4.2 Encrypted Main Memory

- Prior work
 - TRESOR
 - Cryptkeeper
 - Various encrypted, authenticated, and oblivious RAM references
- Keys kept in cache.
- Main memory used as an encrypted swap device
- AESNI
- Security assumptions
 - confidentiality only
 - XTS

4.3 Improved DMA Protections

In order to be of practical use a software cryptoprocessor must allow use of I/O devices, such as those commonly used for storage and networking. These devices typically use DMA to move data between the device and main memory. In conventional operating systems memory for DMA operations is typically allocated from OS memory pools and may even reside on pages with other unrelated data. However, the combination of a cache-resident OS kernel and a security model where devices are treated as malicious brings up several issues that need to be addressed.

4.3.1 Preventing DMA to cleartext data

In a software cryptoprocessor such as vCage only the contents of the cache are in cleartext, while the rest of main memory is encrypted. In order to protect secrets and prevent modification of the system code and data, it is necessary to prevent any DMA operations to the portions of main memory that are cached. This can be achieved by employing DMA protection mechanisms such as Intel VT-d (IOMMU) Protected Memory Regions (PMRs), and Intel VT-d I/O Page Tables (xxx: refs?).

The protected memory regions consist of two pairs of base and bound registers; one only handles memory below 4GB while the other can address all of system memory. This provides considerably less flexibility than using I/O page tables. However, using I/O page tables requires at least a small number of pages from cacheable memory, which is a precious and limited resource. Placing the I/O page tables in uncached (untrusted) memory is not an option since the memory is treated as potentially malicious. In addition, the untrusted memory is marked as uncacheable to ensure it is not fetched into the processor cache, which has a negative impact on the latency of memory accesses to the I/O page tables.

4.3.2 DMA from untrusted memory

Blocking DMA to cacheable memory regions protects the software cryptoprocessor and other cleartext data from being accessed by malicious devices. However, it means that we need to employ some form of bounce buffering (xxx: ref?) so devices are able to perform DMA using uncached memory. This is achieved by allocating a portion of the untrusted memory as a DMA bounce buffer. Unlike the rest of the uncached memory, the contents of this portion of memory are not automatically encrypted by the software cryptoprocessor. It may, however, contain encrypted data if the encryption has been performed by higher levels of the software stack *e.g.* encrypted network traffic.

The Linux kernel provides a standard mechanism for hooking DMA operations by way of a global `dma_ops` variable which points to an instance of `struct dma_map_ops`. We use this mechanism to provide our own set of DMA operations which copy data from system memory to the bounce buffer during `map_page`, `alloc_coherent`, or `sync_single_for_device` operations (xxx: mention ranger operations as well?), and vice-versa during `unmap_page`, `free_coherent`, or `sync_single_for_cpu` operations.

DMA bounce buffering provides the added benefit of preventing time-of-check time-of-use (TOCTOU) attacks. All untrusted data written by an I/O device is copied from the bounce buffer to kernel allocated I/O buffers which reside within the DMA protected region, before any data is made available to other software components. As a result, there is no race condition for a malicious device to exploit.

- Wipe pages before DMA?

4.3.3 Sub-page granularity

Device drivers that allocate memory buffers for use during DMA often do so using the generic `kmalloc` interface, which allocates memory from the size-specific slab caches. In the absence of any interception or modification of DMA transactions, it's possible for a device to access any kernel memory addressable by the DMA engine being used. Even

with IOMMU protection and bounce buffering as outlined above, DMA buffers for multiple devices may end up being allocated from the same physical page, allowing a malicious or errant device to access data being read or written by another device.

In order to provide fine-grained partitioning between devices vCage allocates DMA bounce buffers such that at any point in time the bounce buffer pages associated with different devices are disjoint. In practice this is achieved by allocating the entire page when either the start or end of the DMA buffer only uses some fraction of the page. While this increases the total amount of memory being used by the bounce buffer, since this is only done for *in-flight* DMA operations, the overhead is manageable. (xxx: 'manageable' sounds like a cop out)

4.4 System Attestation

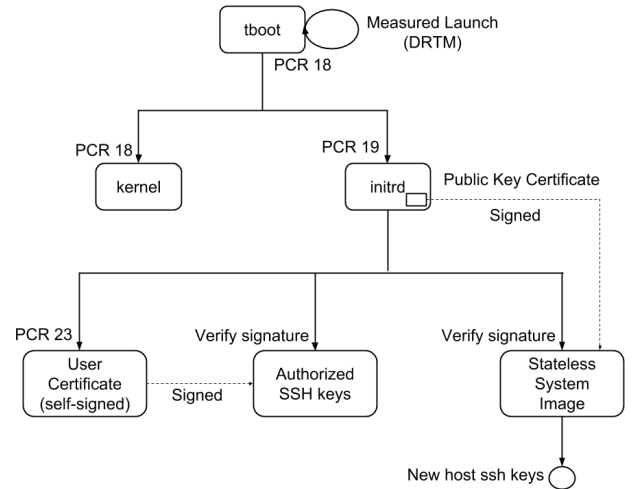


Figure 1: Root of Trust for vCage

- Trusted computing / TPM / TXT
- Chain of measurements
 - tboot → kernel → initrd → keys
 - PCR23
- Stateless
- Keys regenerated on every boot

4.5 System Hardening

Built on a standard Linux kernel, the vCage software cryptoprocessor utilizes various kernel level techniques to increase the security of the system, commonly referred as Linux Kernel Hardening [14]. In this section we describe those techniques in detail, and explain their role in increasing the overall resilience of the system to security threats.

- Minimizing the attack surface: vCage uses a minimized Linux kernel, containing a restricted set of system components and drivers. All interfaces to external devices and filesystems are minimal and carefully selected. In order to further close down any endangering

communication ports, The only communication protocol that is allowed is SSH.

- Enforcing Linux security features: vCage enforces Linux kernel security features including stack canary, Position Independent Executables (PIE), Data Execution Prevention (DEP) and Relocation Read-Only (RELRO) to protect against exploited structures in ELF binaries. Furthermore, it utilizes compilation-time security checks to detect potential string and buffer errors [12].
- Keeping a patched system: Security patches are regularly applied, and custom patches are created when a new vulnerability is disclosed.
- Grsecurity ready: Grsecurity is an extensive security enhancement package to the Linux kernel [1], incorporating numerous changes and touching nearly 2000 files. Some of its features include mitigation of common memory corruption exploits, Mandatory Access Control system and memory protections for userspace applications.

5. EVALUATION

- Real-world use cases
 - commercial product
 - whole-VM protection
 - availability of compatible systems on cloud providers
- More implementation details
- Quantitative experiments

5.1 Confidentiality

- Discussion of correctness
- Eviction counts for various benchmarks

5.2 Performance

- Discuss modifications to Linux kernel / running VMs on KVM
 - Flunky threads
 - Laptop mode
 - scan batch sizes
- Benchmarks
 - Raw memory throughput: 30mb with and w/o privatecore
 - Crypto benchmark: 30mb with and w/o privatecore
 - Webserver benchmark: 30mb with and w/o privatecore
 - Database benchmark: 30mb with and w/o privatecore

6. RELATED WORK

Related work goes here...

Here are a couple of random cites [16, 17] to test the use of our bibliography file.

7. FUTURE WORK

- Authenticated encryption (Steve)
 - GCM performance numbers
 - Merkle-tree rough design, hit estimate
 - McGrew dynamic data sets
- Future performance optimizations?
- Unevictable secrets using address/set mapping?
- Anything else?

8. CONCLUSIONS

Conclusions go here...

9. REFERENCES

- [1] <https://grsecurity.net>.
- [2] A. Boileau. Hit by a Bus: Physical Access Attacks with Firewire. In *RUXCON*, Jan. 2006.
- [3] E. Bugnion, J. M. Anderson, T. C. Mowry, M. Rosenblum, and M. S. Lam. Compiler-directed page coloring for multiprocessors. In *Proceedings of the Seventh International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS VII, pages 244–255, New York, NY, USA, 1996. ACM.
- [4] B. D. Carrier and J. Grand. A hardware-based memory acquisition procedure for digital investigations. *Digital Investigation*, 1(1):50–60, 2004.
- [5] M. Dornseif. Owned by an ipod. In *PacSec*, 2004.
- [6] M. Dornseif. Firewire – all your memory are belong to us. In *CanSecWest*, 2005.
- [7] J. Grand. Patent US7181560 - Method and apparatus for preserving computer memory. US Patent Office, 2007.
- [8] J. A. Halderman, S. D. Schoen, N. Heninger, W. Clarkson, W. Paul, J. A. Calandrino, A. J. Feldman, J. Appelbaum, and E. W. Felten. Lest we remember: cold boot attacks on encryption keys. In *SS'08: Proceedings of the 17th conference on Security symposium*. USENIX Association, July 2008.
- [9] O. Horovitz and S. A. Weis. Physical Privilege Escalation and Mitigation in the x86 World. In *CanSecWest*, 2013.
- [10] Intel Corporation. *Intel 64 and IA-32 Architectures Software Developer's Manual*, June 2013. Order number 325462-047US.
- [11] Intel Corporation. *Intel Virtualization Technology for Directed I/O*, September 2013. Order number D51397-006.
- [12] J. Jelinek. <http://gcc.gnu.org/ml/gcc-patches/2004-09/msg02055.html>, 2004.
- [13] C. Maartmann-Moe. Inception. <http://www.breaknenter.org/projects/inception/>, June 2011.
- [14] T. Merry. <https://www.sans.org/reading-room/whitepapers/linux/linux-kernel-hardening-1294>, 2003.

- [15] C. P. Mozak. Patent US7945050 - Suppressing power supply noise using data scrambling in double data rate memory systems. US Patent Office, 2011.
- [16] T. Müller, F. C. Freiling, and A. Dewald. TRESOR runs encryption securely outside RAM. In *SEC'11: Proceedings of the 20th USENIX Conference on Security*. USENIX Association, Aug. 2011.
- [17] P. Peterson. Cryptkeeper: Improving security with encrypted RAM. In *Technologies for Homeland Security (HST), 2010 IEEE International Conference on*, pages 120–126, 2010.
- [18] N. L. Petroni Jr, T. Fraser, J. Molina, and W. A. Arbaugh. Copilot - A Coprocessor-based Kernel Runtime Integrity Monitor. In *USENIX Security Symposium*, pages 179–194, 2004.
- [19] A. Triulzi. Project Maux Mk. II, I Own the NIC, now I want a shell. *The 8th annual PacSec conference*, 2008.
- [20] C. A. Waldspurger. Memory resource management in VMware ESX server. In *Proceedings of the Fifth Symposium on Operating Systems Design and Implementation*, OSDI '02, pages 181–194, New York, NY, USA, 2002. ACM.
- [21] A. Walters.
http://www.forensicswiki.org/wiki/Volatility_Framework, 2014.
- [22] F. Witherden. Memory Forensics over the IEEE 1394 Interface. *Tech Report*, Sept. 2010.
- [23] X. Zhang, S. Dwarkadas, and K. Shen. Towards practical page coloring-based multicore cache management. In *Proceedings of the 4th ACM European Conference on Computer Systems*, EuroSys '09, pages 89–102, New York, NY, USA, 2009. ACM.